

Quality Attributes

Lecturer: Zhiguo Ding

Phone: 13750985775(660775)

E-mail: dingzhiguo@zjnu.cn

Office: Room 313, Building 21, College of MPI

What are Quality Attributes

- Often know as –ilities
 - Reliability
 - Availability
 - Portability
 - Scalability
 - **Performance** (!)
 - ... but much more than this
- Part of a system's NFRs
 - “how” the system achieves its functional requirements

System Quality Attribute

- ❑ Performance
- ❑ Availability
- ❑ Usability
- ❑ Security

End User's
view

- ❑ Maintainability
- ❑ Portability
- ❑ Reusability
- ❑ Testability

Developer's
view

Business
Community
view

- Time To Market
- Cost and Benefits
- Projected life time
- Targeted Market
- Integration with Legacy System
- Rollout Schedule

**A list of quality attributes exists in
ISO/IEC 9126-2001 Information Technology – Software Product Quality**

Architecture and Quality Attributes

- Achieving quality attributes must be considered throughout design, implementation, and deployment
 - Satisfactory results are a matter of getting 'the big picture'
- Architecture is critical to the realization of many qualities of interest in a system,
 - these qualities should be designed in and can be evaluated at the architectural level.
- Architecture, by itself, is unable to achieve qualities.
 - It provides the foundation for achieving quality, but this foundation will be to no avail if attention is not paid to the details.
- They influence each-other.

Classes of Quality Attributes

- Qualities of the system.
 - availability, modifiability, performance, security, testability, usability, scalability ...
- Business qualities
 - Time to market
 - Cost and benefit
 - Projected lifetime of the system
 - Rollout schedule
- Qualities of the architecture itself
 - Conceptual integrity,
 - Correctness and completeness
 - Buildability

Comments

- Business qualities (such as time to market) are affected by the architecture.
- Iterlinked: Qualities of the architecture itself
 - indirectly affect other qualities, such as modifiability.
 - E.g., conceptual integrity,

System Quality Attributes

- Availability, modifiability, performance, security, testability, usability, scalability ...
- Warning: use **operational** definitions!
 - Architects are often told:
 - “My application must be fast/secure/scale”
 - Far too imprecise to be any use at all
 - Quality attributes (QAs) must be made precise/measurable for a given system design, e.g.
 - *“It must be possible to scale the deployment from an initial 100 geographically dispersed user desktops to 10,000 without an increase in effort/cost for installation and configuration.”*

Quality Attribute Specification

- QA's must be concrete
- But what about testable?
 - Test scalability by installing system on 10K desktops?
- Often careful analysis of a proposed solution is all that is possible
 - “It’s all talk until the code runs”
 - Can you do better?

System Quality Attribute Problems

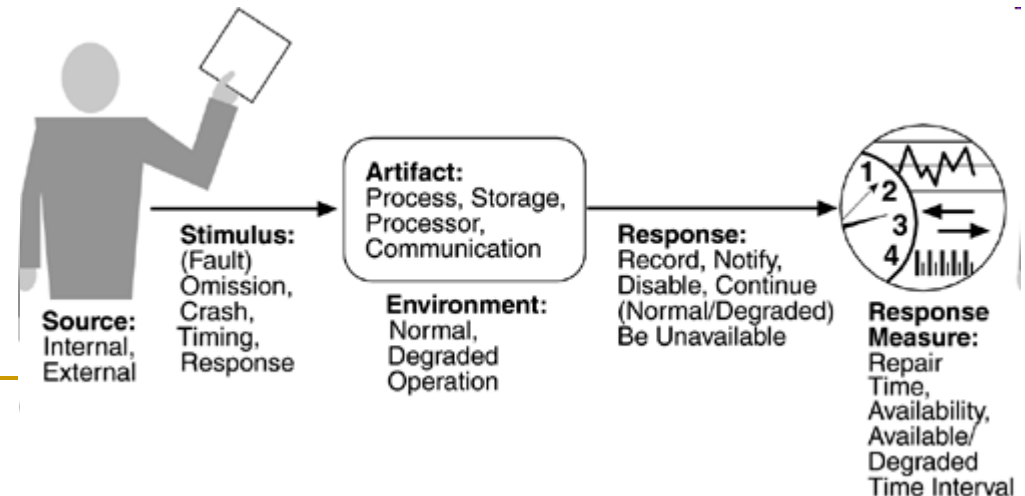
- Definitions are not operational.
- A focus of discussion is often on which quality a particular aspect belongs to.
 - Is a system failure an aspect of availability, an aspect of security, or an aspect of usability? All three attribute communities would claim ownership of a system failure.
- Each attribute community has developed its own vocabulary.
 - performance community → "events"
 - security community → "attacks"
 - the availability community → "failures"
 - the usability community → "user input."
 - All of these may actually refer to the same occurrence,
- Solution: use quality attribute 'scenarios' and unified language

An Analysis Framework for Specifying Quality Attributes

QUALITY ATTRIBUTE SCENARIOS

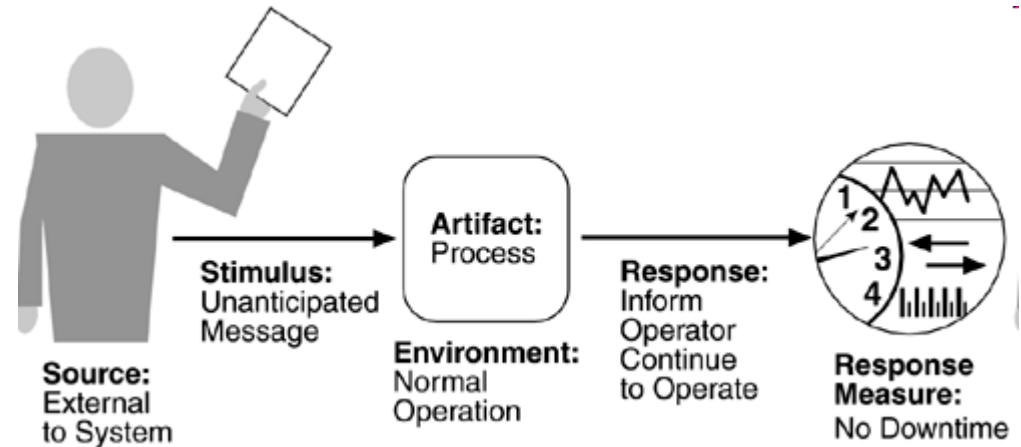
[For each quality-attribute-specific requirement.]

- *Source of stimulus.* This is some entity (a human, a computer system, or any other actuator) that generated the stimulus.
- *Stimulus.* A condition that needs to be considered when it arrives at a system.
- *Environment.* The stimulus occurs within certain conditions. The system may be in an overload condition or may be idle when the stimulus occurs.
- *Artifact.* Some artifact is stimulated. This may be the whole system or some pieces of it.
- *Response.* The activity undertaken after the arrival of the stimulus.
- *Response measure.* When the response occurs, it should be measurable in some fashion so that the requirement can be tested.



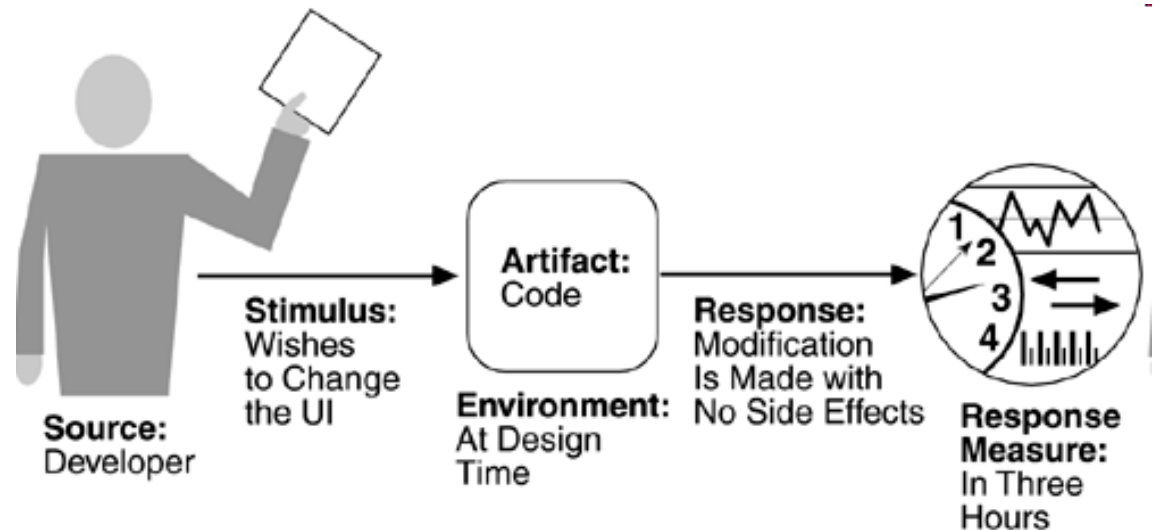
Availability Scenario Example

- System reaction to unanticipated message



Modifiability Scenario Example

- Specifying a modifiability QA requirement



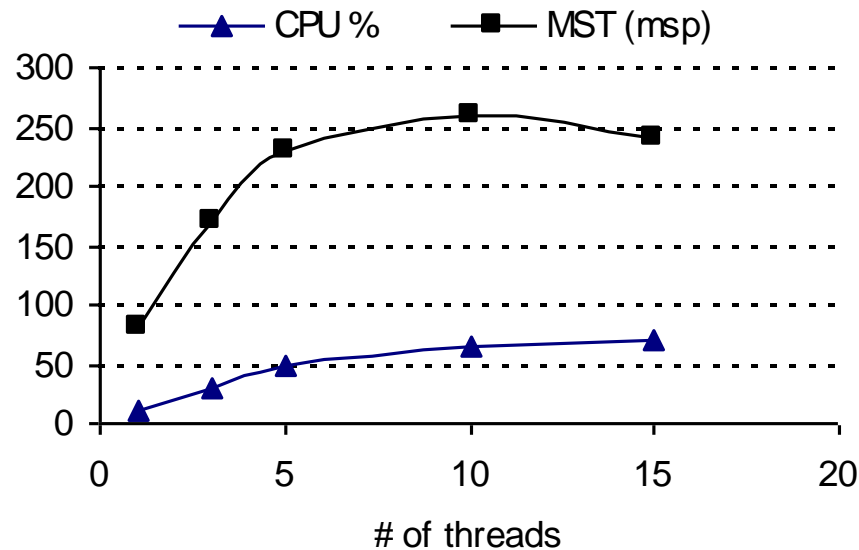
Performance

- Many examples of poor performance in enterprise applications
- Performance requires:
 - Multiple metrics: Throughput, response time, deadlines
 - Average (sustained) vs. peak.
 - Guarantees? Often specified as median and 99% tile.
- Enterprise applications often have strict performance requirements, e.g.
 - 1000 transactions per second
 - 3 second average latency for a request
 - Deadline that must be met

Performance - Throughput

- Measure of the amount of work an application must perform in unit time
 - Transactions per second
 - Messages per minute
- Is required throughput:
 - Average?
 - Peak?
- Many system have low average but high peak throughput requirements

Throughput Example




- Throughput of a message queuing system
 - Messages per second (msp)
 - Maximum sustainable throughput (MST)
- Note throughput changes as number of receiving threads increases

Performance - Response Time

- Measure of the latency an application exhibits in processing a request
- Usually measured in (milli)seconds
- Often an important metric for users
- Is required response time:
 - Guaranteed?
 - Average?
- E.g. 95% of responses in sub-4 seconds, and all within 10 seconds

Response Time

- Example shows response time distribution for a J2EE application



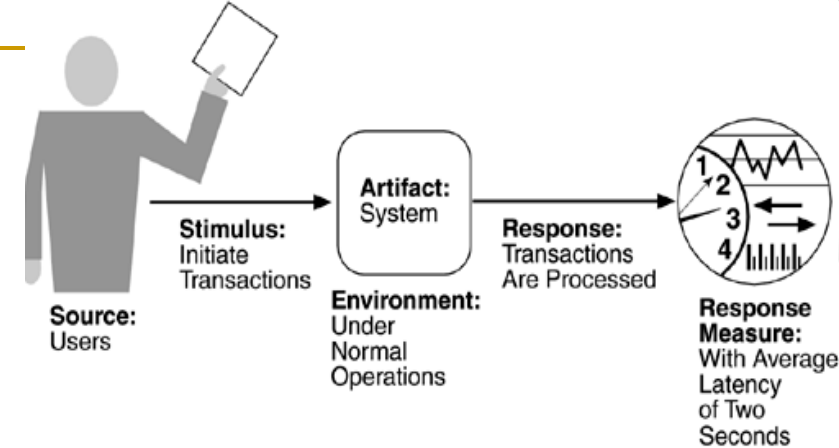
Performance - Deadlines

- ‘something must be completed before some specified time’
 - Payroll system must complete by 2am so that electronic transfers can be sent to bank
 - Weekly accounting run must complete by 6am Monday so that figures are available to management
- Deadlines often associated with batch jobs in IT systems.

Something to watch for ...

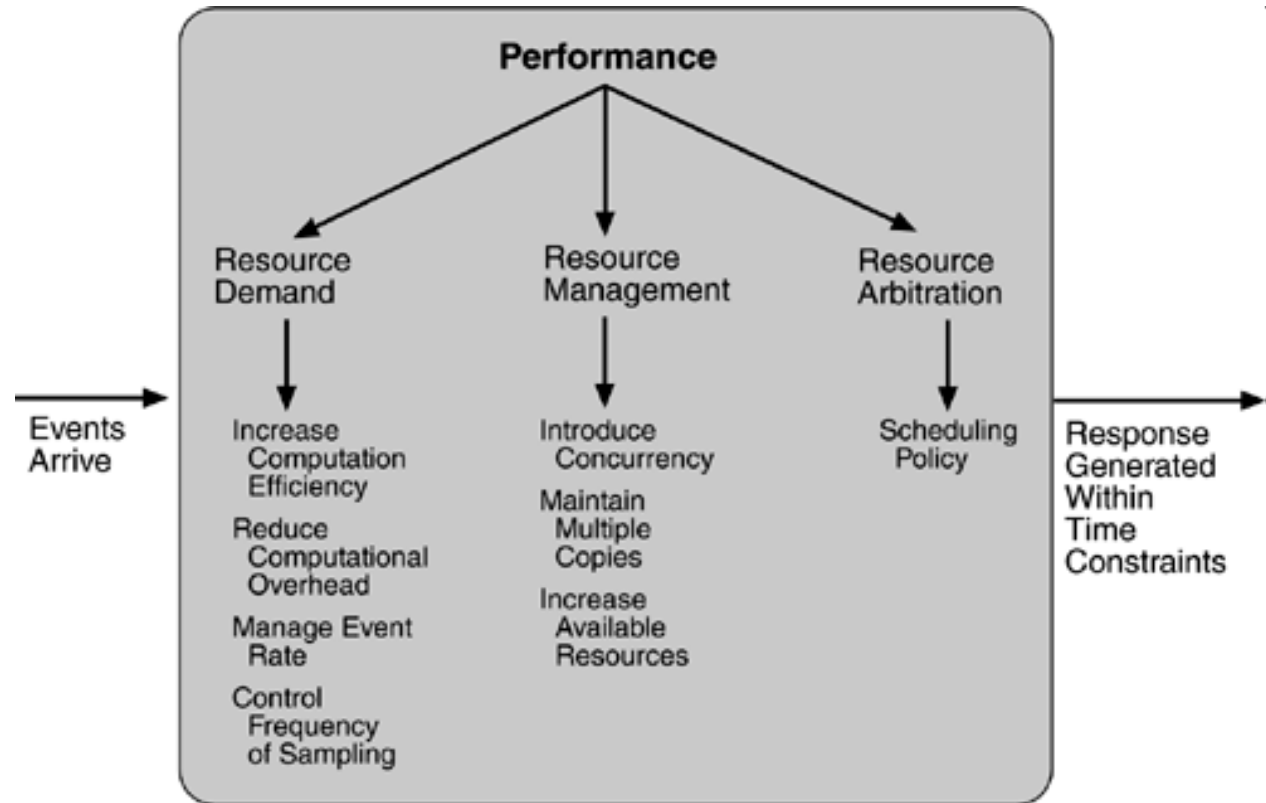
- What is a
 - Transaction?
 - Message?
 - Request?
- All are **application / context specific** measures.
 - System must achieve 100 mps throughput
 - BAD!!
 - System must achieve 100 mps peak throughput for *PaymentReceived* messages
 - GOOD!!!

Performance



Portion of Scenario	Possible Values
<i>Source</i>	One of a number of independent sources, possibly from within system
<i>Stimulus</i>	Periodic events arrive; sporadic events arrive; stochastic events arrive
<i>Artifact</i>	System
<i>Environment</i>	Normal mode; overload mode
<i>Response</i>	Processes stimuli; changes level of service
<i>Response Measure</i>	Latency, deadline, throughput, jitter, miss rate, data loss

Performance Tactics

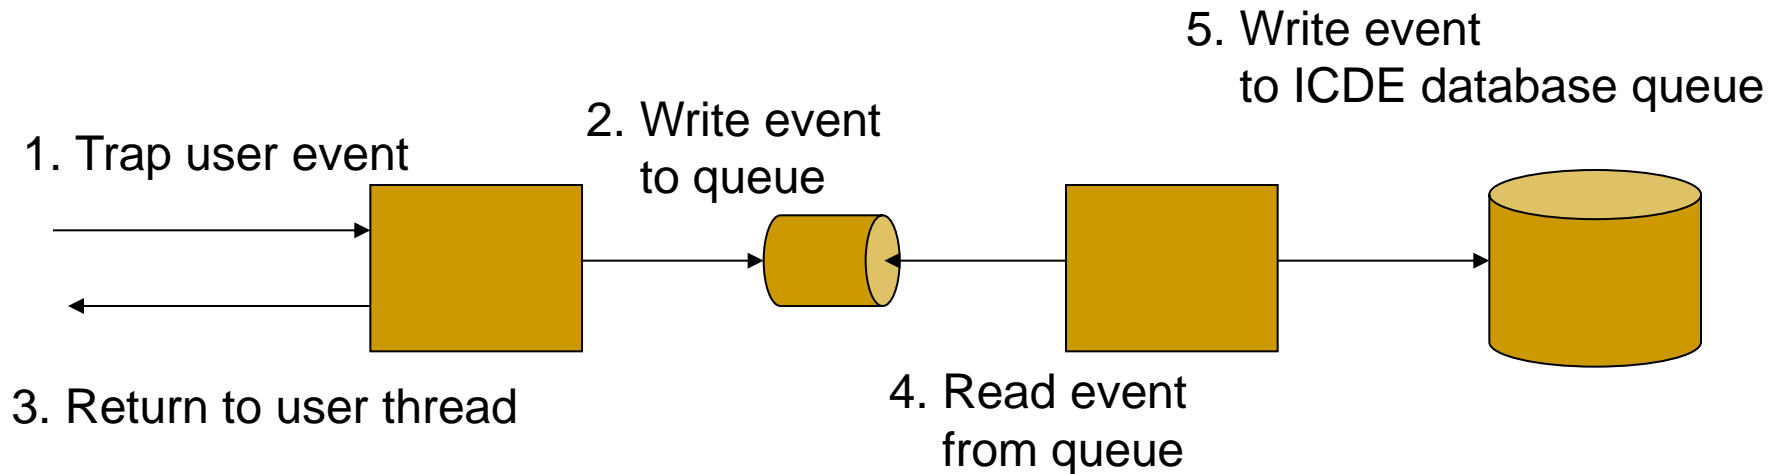


■ Contributors to response time

- Resource consumption
- Blocked time
 - Contention for resources
 - Availability of resources
 - Dependency on other computations

ICDE Performance Issues

- Response time:
 - ❑ Overheads of trapping user events must be imperceptible to ICDE users
- Solution for ICDE client:
 - ❑ Decouple user event capture from storage using a queue



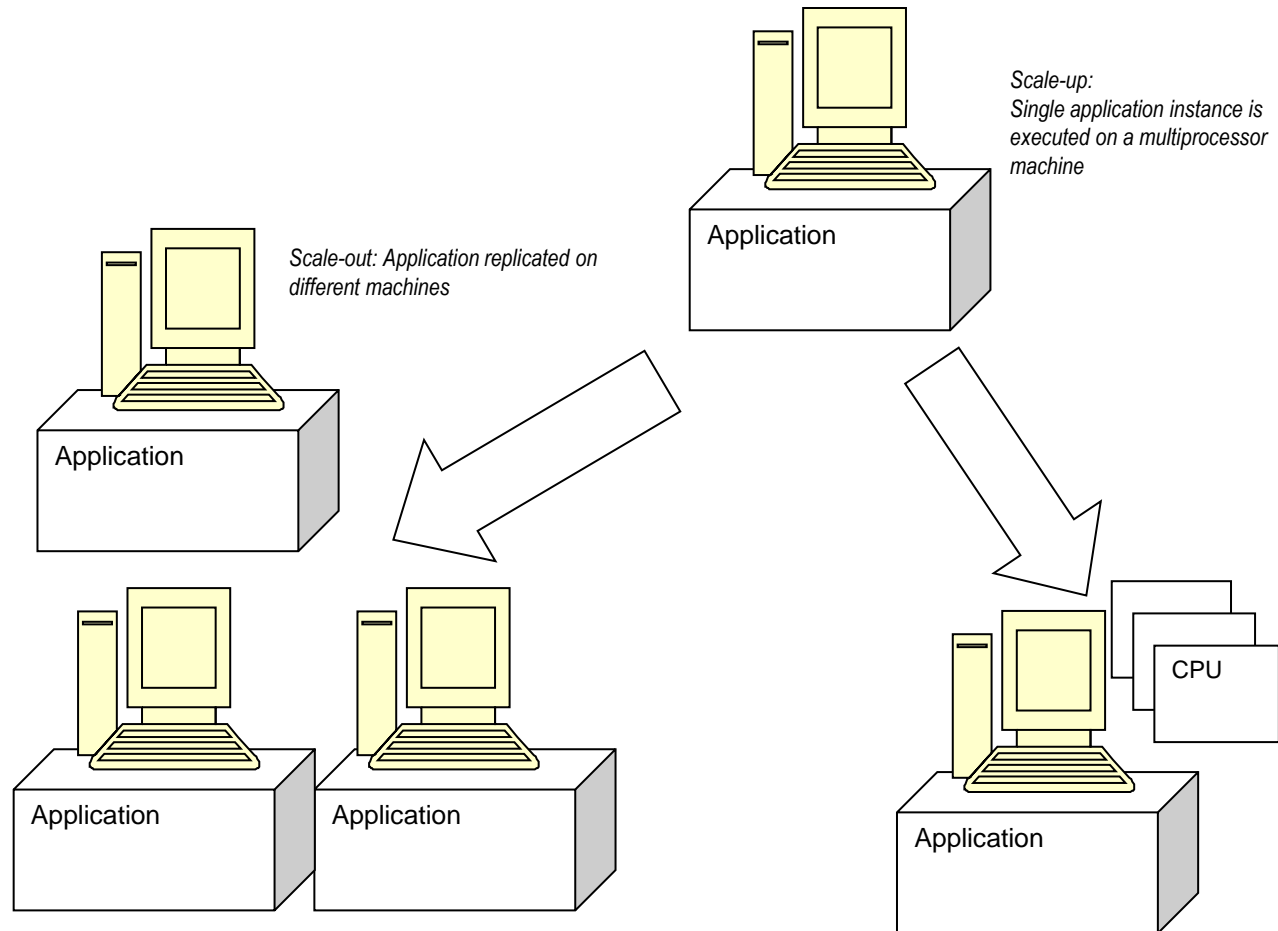
Scalability

- *“How well a solution to some problem will work when the size of the problem increases.”*
- 4 common scalability issues in IT systems:
 - ❑ Request load
 - ❑ Connections
 - ❑ Data size
 - ❑ Deployments

Scalability – Request Load

- How does an 100 tps application behave when simultaneous request load grows? E.g.
 - From 100 to 1000 requests per second?
- Ideal solution, without additional hardware capacity:
 - as the load increases, throughput remains constant (i.e. 100 tps), and response time per request increases only linearly (i.e. 10 seconds).

Scalability – Add more hardware ...



Scalability – the reality

- Adding more hardware should improve performance:
 - scalability must be achieved without modifications to application architecture
- Reality as always is different!
- Applications will exhibit a decrease in throughput and a subsequent exponential increase in response time.
 - increased load causes increased contention for resources such as CPU, network and memory
 - each request consumes some additional resource (buffer space, locks, and so on) in the application, and eventually these are exhausted

Scalability – J2EE example




Figure shows how six different versions of the same application implemented using different JEE application servers perform as their load increases from 100 to 1,000 clients.

Scalability - connections

- What happens if number of simultaneous connections to an application increases
 - ❑ If each connection consumes a resource?
 - ❑ Exceed maximum number of connections?
- ISP example:
 - ❑ Each user connection spawned a new process
 - ❑ Virtual memory on each server exceeded at 2000 users
 - ❑ Needed to support 100Ks of users
 - ❑ Tech crash

Scalability – Data Size

- How does an application behave as the data it processes increases in size?
 - ❑ Chat application sees average message size double?
 - ❑ Database table size grows from 1 million to 20 million rows?
 - ❑ Image analysis algorithm processes images of 100MB instead of 1MB?
- Can application/algorithms scale to handle increased data requirements?

Scalability - Deployment

- How does effort to install/deploy an application increase as installation base grows?
 - Install new users?
 - Install new servers?
- Solutions typically revolve around automatic download/installation
 - E.g. downloading applications from the Internet

Scalability thoughts and ICDE

- Scalability often overlooked.
 - ❑ Major cause of application failure
 - ❑ Hard to predict
 - ❑ Hard to test/validate
 - ❑ Reliance on proven designs and technologies is essential
- For ICDE - application should be capable of handling a peak load of 150 concurrent requests from ICDE clients.
 - ❑ Relatively easy to simulate user load to validate this

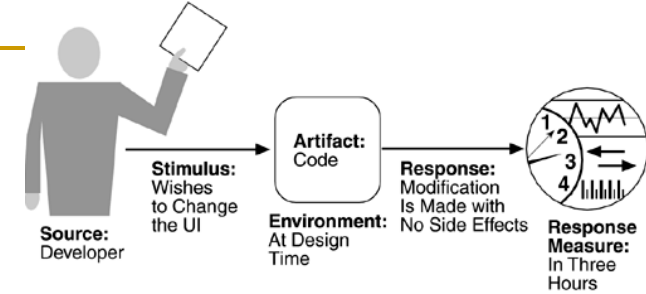
Modifiability

- Modifiability measures how easy it **may** be to change an application to cater for new (non-) functional requirements.
 - ‘**may**’ – nearly always impossible to be certain
 - Must estimate cost/effort
- Modifiability measures are only relevant in the context of a given architectural solution.
 - Components
 - Relationships
 - Responsibilities

Modifiability

- Modifications to a software system during its lifetime are a fact of life.
 - Ideal: modifiable systems that are easier to change/evolve
- Modifiability should be assessed in context of how a system is likely to change
 - No need to facilitate changes that are highly unlikely to occur
 - Over-engineering!
- Impact of designing for modifiability is rarely easy to quantify
- One strategy: Minimizing dependencies
 - Changes isolated to single components likely to be less expensive than those that cause ripple effects across the architecture.

Modifiability



- *Source of stimulus.* Who makes the changes – e.g., developer, a system administrator, or an end user.
- *Stimulus.* What changes? Addition of a function, the modification of an existing function, deletion of a function; changing the qualities of the system
- *Artifact.* Specifies what is to be changed-the functionality of a system, its platform, its user interface, its environment, or another system with which it interoperates.
- *Environment.* When the change can be made-design time, compile time, build time, initiation time, or runtime.
- *Response.* Constraints on the change, how to test and deploy it.
- *Response measure.* Quantitative measure of cost.

Modifiability Scenarios

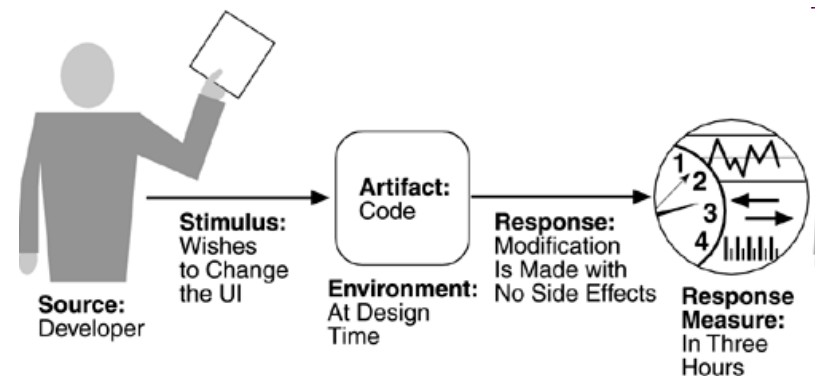
- Provide access to the application through firewalls in addition to existing “behind the firewall” access.
- Incorporate new features for self-service check-out kiosks.
- The COTS speech recognition software vendor goes out of business and we need to replace this component.
- The application needs to be ported from Linux to the Microsoft Windows platform.

Modifiability Analysis


- Impact is rarely easy to quantify
- The best possible is a:
 - Convincing impact analysis of changes needed
 - A demonstration of how the solution can accommodate the modification without change.
- Minimizing dependencies increases modifiability
 - Changes isolated to single components likely to be less expensive than those that cause ripple effects across the architecture.

Modifiability Tactics

- Goals:
 - Reduce the number of modules affected by a change
 - → localize modifications
 - Limited modifications of these modules
 - → prevent ripple effects
 - Control deployment time and cost
 - → defer binding time



Modifiability Tactics



■ Goals:

- ❑ Reduce the number of modules affected by a change
 - ➔ localize modifications
- ❑ Limited modifications of these modules
 - ➔ prevent ripple effects
- ❑ Control deployment time and cost
 - ➔ defer binding time

■ Techniques

- ❑ *Maintain semantic code coherence*
 - Coupling & cohesion metrics
- ❑ *Abstract common services*
- ❑ *Anticipate expected changes*
- ❑ *Generalize the module*

Modifiability Tactics

■ Goals:

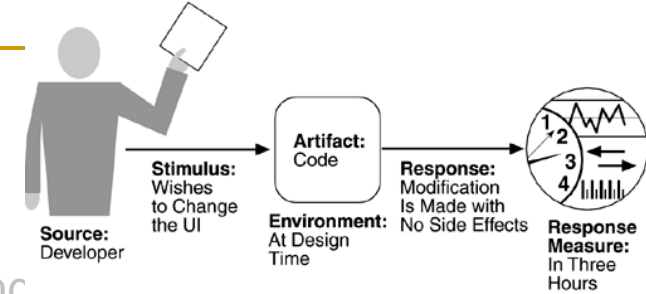
- Reduce the number of modules affected by a change
 - → localize modifications
- **Limited modifications of these modules**
 - → **prevent ripple effects**
- Control deployment time and cost
 - → defer binding time

A taxonomy of dependencies between modules

- *Syntax of data and service invocation*
- *Semantics of data and service*
- *Sequence of data and control*
- *Identity of interfaces*
- *Location of called service*
- *QoS provided*

Tactics:

- *Hide information*
- *Maintain existing interfaces*
- *Restrict communication paths*
- *Use an intermediary (proxy)*



Modifiability Tactics

■ Goals:

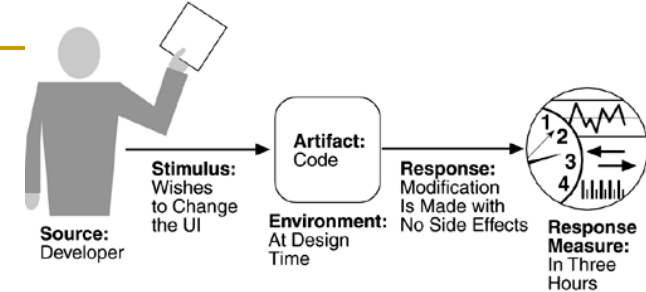
- ❑ Reduce the number of modules affected by a change
 - → localize modifications
- ❑ Limited modifications of these modules
 - → prevent ripple effects
- ❑ Control deployment time and cost
 - → defer binding time

■ Issues:

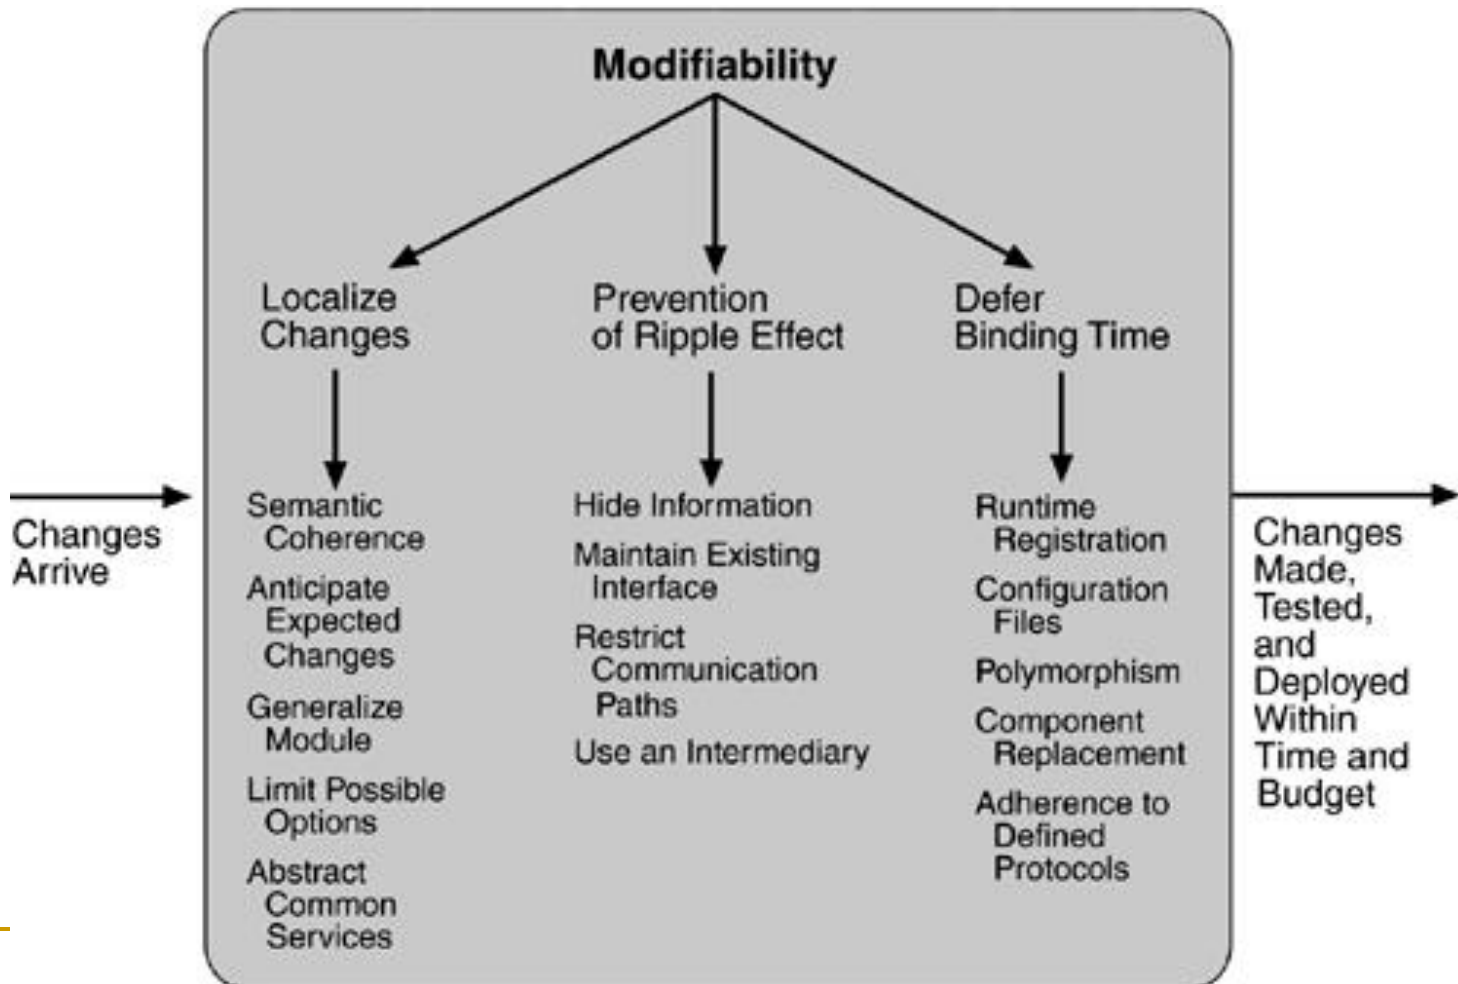
- ❑ Reduce time to deploy
- ❑ Allow non-programmers to make changes

■ *Tactics*

- ❑ Runtime registration
- ❑ Configuration files
- ❑ Dynamically loaded code



Modifiability Tactics Summary



Modifiability for ICDE

- The range of events trapped and stored by the ICDE client to be expanded.
- Third party tools to communicate new message types.
- Change database technology used
- Change server technology used

Security

- Difficult, specialized quality attribute:
 - Lots of technology available
 - Requires deep knowledge of approaches and solutions
- Security is a multi-faceted quality ...

Security

- **Authentication:** Applications can verify the identity of their users and other applications with which they communicate.
- **Authorization:** Authenticated users and applications have defined access rights to the resources of the system.
- **Encryption:** The messages sent to/from the application are encrypted.
- **Integrity:** This ensures the contents of a message are not altered in transit.
- **Non-repudiation:** The sender of a message has proof of delivery and the receiver is assured of the sender's identity. This means neither can subsequently refute their participation in the message exchange.

Security Approaches

- SSL(Security Socket Layer)
- PKI(Public Key Infrastructure)
- Web Services security
- JAAS
 - Java Authentication and Authorization Service
- Operating system security
- Database security
- Etc.

ICDE Security Requirements

- Authentication of ICDE users and third party ICDE tools to ICDE server
- Encryption of data to ICDE server from 3rd party tools/users executing remotely over an insecure network

Availability

- Key requirement for most IT applications
- Measured by the proportion of the required time it is useable. E.g.
 - E.g.
 - 100% available during business hours
 - No more than 2 hours scheduled downtime per week
 - 24x7x52 (100% availability)
 - $MTTF / (MTTF + MTTR)$
 - Note: scheduled downtime is excluded
- Related to an application's reliability
 - Unreliable applications suffer poor availability

Availability


- Period of loss of availability determined by:
 - Time to detect failure
 - Time to correct failure
 - Time to restart application
- Strategies for high availability:
 - Eliminate single points of failure
 - Replication and failover
 - Automatic detection and restart
- Distinguish between faults and failures
- Recoverability (e.g. a database)
 - the capability to reestablish performance levels and recover affected data after an application or system failure

Availability Scenarios

- *Source of stimulus.* Internal or external indications of faults or failure since the desired system response may be different.
- *Stimulus.* A fault of one of the following classes occurs.
 - ❑ omission. A component fails to respond to an input
 - ❑ crash. The component repeatedly suffers omission faults.
 - ❑ timing. A component responds but the response is early or late.
 - ❑ response (byzantine). A component responds with an incorrect value.
- *Artifact.* Specifies the resource that is required to be available,
- *Environment.* The state of the system when the fault or failure occurs may also affect the desired system response
- *Response.* Possible reactions to a system failure:
 - ❑ logging the failure, notifying selected users or other systems, switching to a degraded mode with either less capacity or less function, shutting down external systems, or becoming unavailable during repair.
- *Response measure.* Metric of success: e.g., availability percentage, or time to repair, etc.

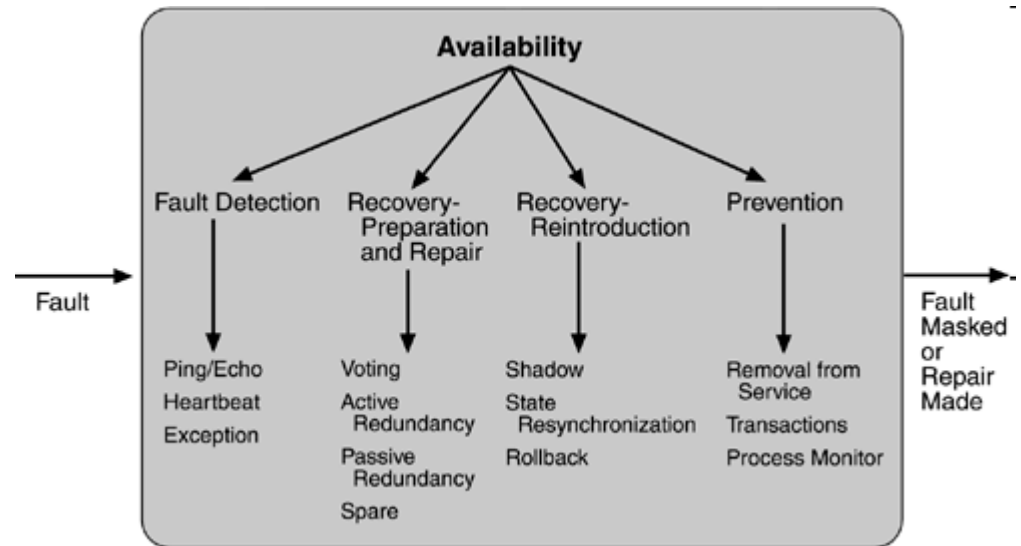
Designing for Availability

■ Faults vs. Failures



■ Tactics

- ❑ Fault detection
- ❑ Fault recovery
- ❑ Fault prevention



Tactics for Availability

■ Fault detection

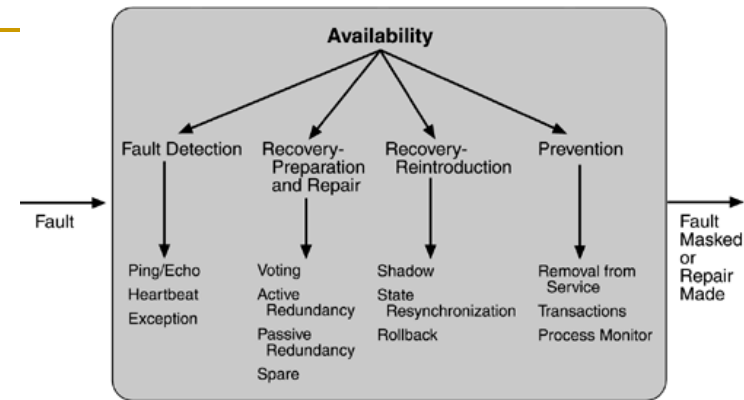
- ❑ Ping/echo;
- ❑ Heartbeat;
- ❑ Exceptions

■ Fault recovery

- ❑ Mostly redundancy based
 - [byzantine faults] Voting: multiple processes working in parallel.
 - [crash, timing] Active redundancy – hot restart
 - [crash] Passive redundancy (warm restart), spare.
- ❑ Reintroduction: shadow operation, resynchronization, checkpoint/rollback

■ Fault prevention

- ❑ Removal from service; Transactions



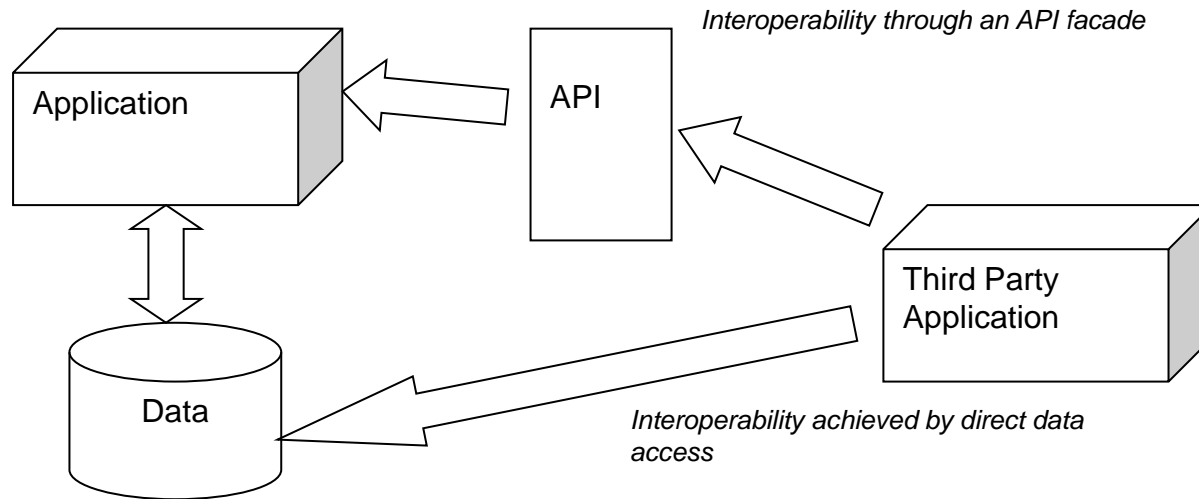
Availability for ICDE

- Achieve 100% availability during business hours
- Plenty of scope for downtime for system upgrade, backup and maintenance.
- Include mechanisms for component replication and failover

Integration

- Ease with which an application can be incorporated into a broader application context
 - Use component in ways that the designer did not originally anticipate
- Typically achieved by:
 - Programmatic APIs
 - Data integration

Integration Strategies



- Data – expose application data for access by other components
- API – offers services to read/write application data through an abstracted interface
- Each has strengths and weaknesses ...

ICDE Integration Needs

- Revolve around the need to support third party analysis tools.
- Well-defined and understood mechanism for third party tools to access data in the ICDE data store.

Other Quality Attributes

■ Portability

- Can an application be easily executed on a different software/hardware platform to the one it has been developed for?

■ Testability

- How easy or difficult is an application to test?

■ Supportability

- How easy an application is to support once it is deployed?

Other Quality Attributes

■ Portability

- Can an application be easily executed on a different software/hardware platform to the one it has been developed for?

■ Testability

- How easy or difficult is an application to test?

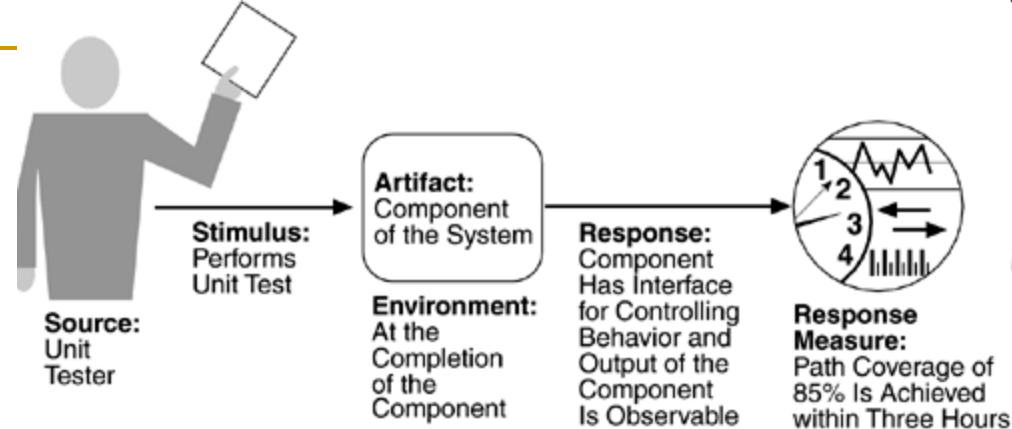
■ Supportability

- How easy an application is to support once it is deployed?

Testability

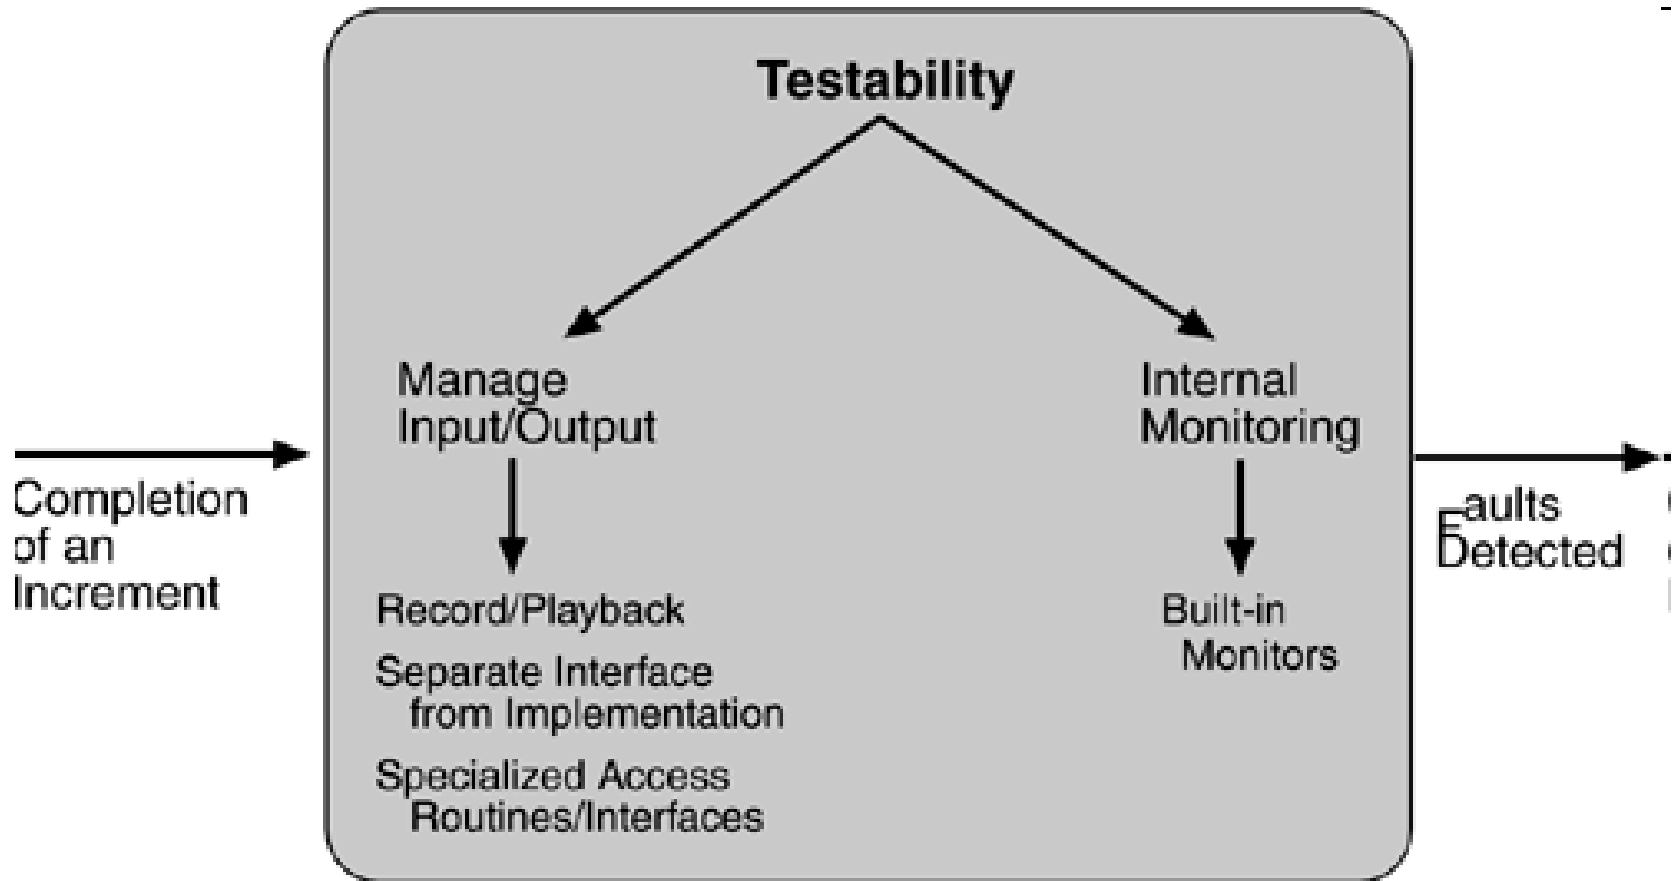
- Estimate: 40% of development cost goes to testing
- Testability: assuming that the software has at least one fault, the probability that this will be detected in the next testing round
- Need a system that is *controllable* and *observable*
- **Testing harness**: control internal state of components,, pass inputs to the system, observe output

Testability



Portion of Scenario	Possible Values
Source	Unit test developer; or Increment integrator; or System verifier; or Client acceptance tester; or System user
Stimulus	Analysis, architecture, design, class, subsystem integration completed; system delivered
Artifact	Piece of design, piece of code, complete application
Environment	At design time, at development time, at compile time, at deployment time
Response	Provides access to state values; provides computed values; prepares test environment
Response Measure	Percent executable statements executed Probability of failure if fault exists Time to perform tests Length of longest dependency chain in a test Length of time to prepare test environment

Testability Tactics



Other Quality Attributes

■ Portability

- Can an application be easily executed on a different software/hardware platform to the one it has been developed for?

■ Testability

- How easy or difficult is an application to test?

■ Supportability

- How easy an application is to support once it is deployed?

Concern:

- Q: Are functionality and quality attributes orthogonal?
- Q: How to generate meaningful scenarios in practice?
- Q: Completeness?

Design Trade-offs

- QAs are rarely orthogonal
 - They interact, affect each other
 - highly secure system may be difficult to integrate
 - highly available application may trade-off lower performance for greater availability
 - high performance application may be tied to a given platform, and hence not be easily portable
- Architects must create solutions that makes sensible design compromises
 - not possible to fully satisfy all competing requirements
 - Must satisfy all stakeholder needs
 - This is the difficult bit!

Summary

- QAs are part of an application's non-functional requirements
- Many QAs
- Architect must decide which are important for a given application
 - Understand implications for application
 - Understand competing requirements and trade-offs

Selected Further Reading

- Ian Gorton, **Essential Software Architecture** (2nd Edition) , Springer; 2011.
- L. Chung, B. Nixon, E. Yu, J. Mylopoulos, (Editors). Non-Functional Requirements in Software Engineering Series: The Kluwer International Series in Software Engineering. Vol. 5, Kluwer Academic Publishers. 1999.
- J. Ramachandran. Designing Security Architecture Solutions. Wiley & Sons, 2002.
- I.Gorton, L. Zhu. *Tool Support for Just-in-Time Architecture Reconstruction and Evaluation: An Experience Report*. International Conference on Software Engineering (ICSE) 2005, St Louis, USA, ACM Press

Thank you !
